

```
1 # Install CRAN packages if missing
2 packages <- c("dplyr", "ggplot2", "lme4", "emmeans", "car", "ARTool", "rstatix", "effectsize", "effsize")
3
4 installed <- rownames(installed.packages())
5 for (p in packages) {
6   if (!(p %in% installed)) {
7     install.packages(p, dependencies = TRUE)
8   }
9 }
10
11 # Load them after installation
12 lapply(packages, library, character.only = TRUE)
13
```



Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Loading required package: Matrix

▼ Data Import & Cleaning

Welcome to emmeans

Caution: You lose important information if you filter this package's results.

See '? untidy'

We load the experimental run table (`run_table.csv`), inspect its structure, and remove missing values.

We also check the number of runs per treatment combination to ensure balanced data.

⚠ In a later step, we must report failed runs (e.g., "X out of N failed due to Y") instead of silently removing them.

Attaching package: 'car'

```
1 # Load the CSV
2 df <- read.csv("run_table.csv")
3
4 # Check structure
5 str(df)
6
7 # Basic cleaning: remove NAs
8 df <- na.omit(df)
9
10 # Summary of missing values
11 colSums(is.na(df))
12
13 # Check number of runs per treatment combination
14 df %>% count(gc_strategy, workload, jdk)
15
16 # Count failed runs before NA removal
17 failed_runs <- sum(is.na(df))
18 total_runs <- nrow(df) + failed_runs
19 cat("Failed runs:", failed_runs, "out of", total_runs, "\n")
20
```

```
1. 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
2. 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
3. 'lme4' · 'Matrix' · 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
4. 'emmeans' · 'lme4' · 'Matrix' · 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
5. 'car' · 'carData' · 'emmeans' · 'lme4' · 'Matrix' · 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
6. 'ARTool' · 'car' · 'carData' · 'emmeans' · 'lme4' · 'Matrix' · 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
7. 'rstatix' · 'ARTool' · 'car' · 'carData' · 'emmeans' · 'lme4' · 'Matrix' · 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
8. 'effectsize' · 'rstatix' · 'ARTool' · 'car' · 'carData' · 'emmeans' · 'lme4' · 'Matrix' · 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
9. 'effsize' · 'effectsize' · 'rstatix' · 'ARTool' · 'car' · 'carData' · 'emmeans' · 'lme4' · 'Matrix' · 'ggplot2' · 'dplyr' · 'stats' · 'graphics' · 'grDevices' · 'utils' · 'datasets' · 'methods' · 'base'
```

```
'data.frame': 216 obs. of 12 variables:
 $ X_run_id      : chr "run_0_repetition_0" "run_1_repetition_0" "run_2_repetition_0" "run_3_repetition_0" ...
 $ X_done        : chr "DONE" "DONE" "DONE" "DONE" ...
 $ subject       : chr "dacapo" "dacapo" "dacapo" "dacapo" ...
 $ gc_strategy   : chr "SerialGC" "SerialGC" "SerialGC" "SerialGC" ...
 $ workload      : chr "light" "light" "medium" "medium" ...
 $ jdk           : chr "oraclejdk11" "openjdk17" "oraclejdk11" "openjdk17" ...
 $ energy_joules : num 3.58 2.57 7.66 8.98 9.51 ...
 $ execution_time: num 0.327 0.338 0.649 0.71 1.036 ...
 $ power_watts   : num 10.96 7.61 11.81 12.64 9.18 ...
 $ exit_code     : int 0 0 0 0 0 0 0 0 ...
 $ X             : int 1 2 3 4 5 6 7 8 9 10 ...
 $ X.1           : chr "/Users/rahilsharma/Desktop/Green-Lab/Mock-Server/gc_energy_experiment/experiments/gc_energy_experiment/run_0_repetition_0" "/Users/rahilsharma/Desktop/Green-Lab/Mock-Server/gc_energy_experiment/experiments/gc_energy_experiment/run_1_repetition_0" ...

X_run_id: 0 X_done: 0 subject: 0 gc_strategy: 0 workload: 0 jdk: 0 energy_joules: 0 execution_time: 0 power_watts: 0 exit_code: 0 X: 0 X.1: 0

A data.frame: 18 x 4
  gc_strategy workload      jdk      n
    <chr>      <chr>    <chr> <int>
1 G1GC        heavy   openjdk17    12
2 G1GC        heavy   oraclejdk11   12
3 G1GC        light   openjdk17    12
4 G1GC        light   oraclejdk11   12
5 G1GC        medium  openjdk17    12
6 G1GC        medium  oraclejdk11   12
7 ParallelGC   heavy   openjdk17    12
8 ParallelGC   heavy   oraclejdk11   12
9 ParallelGC   light   openjdk17    12
10 ParallelGC  light   oraclejdk11   12
11 ParallelGC  medium  openjdk17    12
12 ParallelGC  medium  oraclejdk11   12
13 SerialGC     heavy   openjdk17    12
14 SerialGC     heavy   oraclejdk11   12
15 SerialGC     light   openjdk17    12
16 SerialGC     light   oraclejdk11   12
17 SerialGC     medium  openjdk17    12
18 SerialGC     medium  oraclejdk11   12

Failed runs: 0 out of 216
```

▼ Descriptive Statistics & Visualization

We compute the mean, standard deviation, and 95% confidence intervals for energy consumption across GC, workload, and JDK combinations.

We also visualize distributions using boxplots and interaction plots to identify trends and possible outliers.

```
1 # Descriptive statistics
2 df %>%
3   group_by(gc_strategy, workload, jdk) %>%
4   summarise(
5     mean_energy = mean(energy_joules),
6     sd_energy   = sd(energy_joules),
7     ci_low      = mean_energy - qt(0.975, n()-1) * sd_energy/sqrt(n()),
8     ci_high     = mean_energy + qt(0.975, n()-1) * sd_energy/sqrt(n()),
9     .groups = "drop"
10  )
11
12 # Boxplots
13 ggplot(df, aes(x = gc_strategy, y = energy_joules, fill = gc_strategy)) +
14   geom_boxplot() +
15   facet_wrap(~workload) +
16   theme_minimal()
17
18 # Interaction plot
19 interaction.plot(df$gc_strategy, df$workload, df$energy_joules)
20
21
22 # Identify outliers in energy consumption
23 outliers <- boxplot.stats(df$energy_joules)$out
24 outliers
25
26 # Inspect which runs are outliers
27 df %>% filter(energy_joules %in% outliers)
28
29 outliers <- boxplot.stats(df$energy_joules)$out
30 print(outliers)
31
```



A tibble: 18 × 7

gc_strategy	workload	jdk	mean_energy	sd_energy	ci_low	ci_high
<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
G1GC	heavy	openjdk17	15.370162	4.3996477	12.574760	18.165565
G1GC	heavy	oraclejdk11	11.477760	4.9121901	8.356703	14.598816
G1GC	light	openjdk17	4.190206	1.2493240	3.396423	4.983989
G1GC	light	oraclejdk11	3.982909	0.9971837	3.349328	4.616489
G1GC	medium	openjdk17	9.603073	3.0567949	7.660878	11.545268
G1GC	medium	oraclejdk11	8.954945	3.4731396	6.748217	11.161672
ParallelGC	heavy	openjdk17	14.674603	2.7285167	12.940986	16.408220
ParallelGC	heavy	oraclejdk11	13.853493	2.9230599	11.996269	15.710717
ParallelGC	light	openjdk17	4.096408	0.6394977	3.690090	4.502725
ParallelGC	light	oraclejdk11	4.545249	0.9721383	3.927582	5.162916
ParallelGC	medium	openjdk17	10.468336	2.7424103	8.725892	12.210781
ParallelGC	medium	oraclejdk11	10.718180	2.1917966	9.325579	12.110781
SerialGC	heavy	openjdk17	11.125027	1.9343769	9.895982	12.354071
SerialGC	heavy	oraclejdk11	12.203132	2.0193407	10.920105	13.486160

Boxplots showed a few potential outliers, but none were flagged as extreme by the boxplots. Therefore, no runs were excluded.

▼ Data Preparation

SerialGC	light	oraclejdk11	3.209871	0.6046521	2.825693	3.594048
SerialGC	medium	openjdk17	8.220377	0.9042711	7.645830	8.794923
SerialGC	medium	oraclejdk11	7.952360	0.8858042	7.389546	8.515173

We ensure that categorical variables (GC strategy, workload, JDK, subject) are properly treated as factors for analysis.

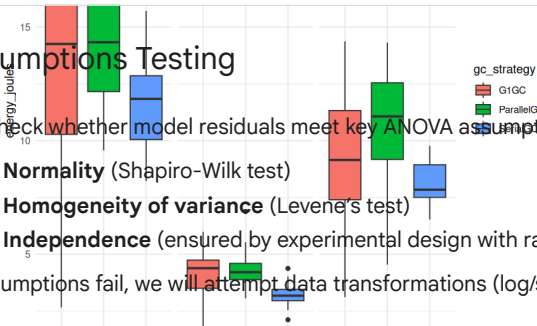
```
1 df$gc_strategy <- factor(df$gc_strategy)
2 df$workload <- factor(df$workload)
3 df$jdk <- factor(df$jdk)
4 df$subject <- factor(df$subject)
```

▼ Assumptions Testing

We check whether model residuals meet key ANOVA assumptions:

- **Normality** (Shapiro-Wilk test)
- **Homogeneity of variance** (Levene's test)
- **Independence** (ensured by experimental design with randomization and cooldowns).

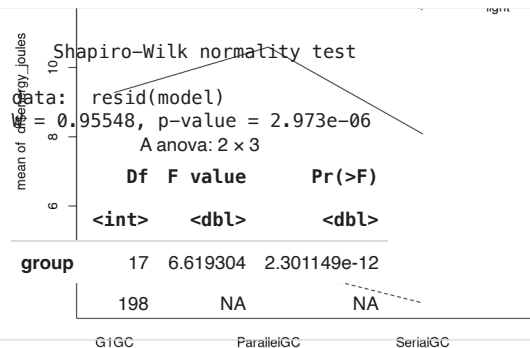
If assumptions fail, we will attempt data transformations (log/sqrt) before switching to non-parametric methods.



```

1 # Fit linear model for residuals check
2 model <- lm(energy_joules ~ gc_strategy * workload * jdk, data=df)
3
4 # Normality test on residuals
5 shapiro.test(resid(model))
6
7 # Homogeneity of variance (Levene's test)
8 leveneTest(energy_joules ~ gc_strategy * workload * jdk, data=df)
9
10 # Independence: assumed from randomization protocol
11

```



Both tests showed significant violations ($p < 0.001$), indicating the ANOVA assumptions are not satisfied. Therefore, we attempted a log-transformation before considering non-parametric methods.

✓ Data Transformation Attempt (Log)

Since the Shapiro-Wilk and Levene's tests showed that assumptions of normality and equal variance are violated, we attempt a simple log-transformation of the energy values.

The results still indicated strong violations of normality ($p < 1e-9$) and homogeneity ($p \approx 0.006$).

Thus, transformation did not resolve assumption issues, and we proceed with ART (non-parametric factorial ANOVA).

```

1 # Apply log transformation to the dependent variable
2 df$log_energy <- log(df$energy_joules)
3
4 # Check assumptions again with transformed data
5 log_model <- lm(log_energy ~ gc_strategy * workload * jdk, data=df)
6
7 # Shapiro-Wilk test for normality on residuals
8 shapiro.test(resid(log_model))
9
10 # Levene's test for homogeneity of variance
11 leveneTest(log_energy ~ gc_strategy * workload * jdk, data=df)
12

```

```

Shapiro-Wilk normality test

data:  resid(log_model)
W = 0.91183, p-value = 5.084e-10
A anova: 2 x 3

  Df F value    Pr(>F)
  <int>  <dbl>    <dbl>
group    17  2.176071 0.005905881
        198      NA      NA

```

Log-transformation did not resolve assumption violations (normality and homogeneity still violated). Therefore, we proceed with non-parametric ART for hypothesis testing.

✓ Hypothesis Testing (Parametric, for comparison)

Although assumptions were not satisfied, we still ran a mixed-effects ANOVA to provide a comparison.

Results showed workload as the dominant factor ($F \approx 265$, $p < 1e-50$) and GC strategy also significant ($F \approx 13.5$, $p < 0.001$).

However, due to violated assumptions, these results should be interpreted with caution and are presented as supplementary.

```

1 # Mixed-effects ANOVA: subject as random effect
2 anova_model <- lmer(energy_joules ~ gc_strategy * workload * jdk + (1|subject), data=df)
3
4 anova(anova_model)
5
6 # Post-hoc Tukey HSD
7 emmeans(anova_model, pairwise ~ gc_strategy | workload)
8
9
10 # Example: adjust p-values from Tukey comparisons
11 tukey_results <- emmeans(anova_model, pairwise ~ gc_strategy | workload)
12 tukey_df <- as.data.frame(tukey_results$contrasts)
13
14 # Apply Benjamini-Hochberg correction
15 tukey_df$p.adjusted <- p.adjust(tukey_df$p.value, method = "BH")
16 tukey_df
17
18

```




```
A anova: 7 x 4
      npar      Sum Sq      Mean Sq      F value
      <int>      <dbl>      <dbl>      <dbl>
gc_strategy 2 157.96009  78.980043 13.4734016
workload     2 3111.92855 1555.964274 265.4358085
jdk          1  10.87962  10.879618  1.8559811
gc_strategy:workload 4  18.80326   4.700816  0.8019239
gc_strategy:jdk     2  35.61363  17.806817  3.0377091
workload:jdk        2  16.57793   8.288965  1.4140351
gc_strategy:workload:jdk 4  43.64785  10.911962  1.8614987
```

NOTE: Results may be misleading due to involvement in interactions

```
$emmeans
workload = heavy:
gc_strategy emmean SE df lower.CL upper.CL
G1GC        13.42 0.52 61   12.38   14.46
ParallelGC   14.26 0.52 61   13.22   15.30
SerialGC     11.66 0.52 61   10.62   12.70
```

✓ Hypothesis Testing (Non-Parametric ART)

```
workload = light:
gc_strategy emmean SE df lower.CL upper.CL
G1GC        4.09 0.52 61    3.05    5.13
ParallelGC   4.32 0.52 61    3.28    5.36
SerialGC     5.20 0.52 61    4.16    6.24
```

Given that both raw and log-transformed data violated assumptions, we applied the Aligned Rank Transform (ART) for factorial ANOVA.

- workload = medium:
- Results confirmed:
- **GC strategy**: significant ($F \approx 24.3$, $p = 0.000159$)
 - **Workload**: dominant effect ($F \approx 265$, $p < 1e-56$)
 - **JDK**: not significant ($p \approx 0.09$)
 - **GC x workload interaction**: keyword - **yes** (borderline ($p \approx 0.058$))
- Confidence level used: 0.95

We therefore base our main conclusions on these ART results.

```
$contrasts
workload = heavv:

1 art_model <- art(energy_joules ~ gc_strategy * workload * jdk + (1|subject), data=df)
2 anova(art_model)

ParallelGC - SerialGC    2.600 0.699 195    3.720 0.0008

workload = light:
contrast      estimate      SE df t.ratio p.value
G1GC - ParallelGC   -0.234 0.699 195   -0.335 0.9400
G1GC - SerialGC      0.887 0.699 195    1.269 0.4144
ParallelGC - SerialGC  1.121 0.699 195    1.604 0.2463

workload = medium:
contrast      estimate      SE df t.ratio p.value
G1GC - ParallelGC   -1.314 0.699 195   -1.880 0.1471
G1GC - SerialGC      1.193 0.699 195    1.706 0.2054
ParallelGC - SerialGC  2.507 0.699 195    3.587 0.0012
```

Results are summed over the levels of all

Results are averaged over the levels of: jdk
Degrees-of-freedom method: kenward-rooper
P value adjustment: tukey method for comparing a family of estimates
NOTE: Results may be misleading due to involvement in interactions

	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	
gc_strategy	gc_strategy	24.341860	2	195	3.654175e-10	
workload	workload	265.174416	2	195	2.371648e-56	p.adjusted
jdk	jdk	2.881885	1	195	9.117637e-02	
gc_strategy:workload	gc_strategy:workload	2.327720	4	195	5.765203e-02	
G1GC - SerialGC	heavy	1.7598815	0.6989232	195	2.5179900	0.0335814422
gc_strategy:jdk	gc_strategy:jdk	4.368306	2	195	1.393604e-02	0.005518166
workload:jdk	workload:jdk	1.715447	2	195	1.825865e-01	
G1GC - ParallelGC	light	-0.2342712	0.6989232	195	-0.3351888	0.9399692137
gc_strategy:workload:jdk	gc_strategy:workload:jdk	2.024419	4	195	9.251673e-02	0.510107108

6	ParallelGC - SerialGC	light	1.1211747	0.6989232	195	1.6041459	0.2463069145	0.369460372
2	ParallelGC - ParallelGC	medium	-1.3142492	0.6989232	195	-1.8803916	0.1470940800	0.330961680
8	G1GC - SerialGC	medium	1.1926408	0.6989232	195	1.7063976	0.2053624344	0.369460372

We compute Cohen's d for pairwise comparisons of GC strategies.
This complements p-values by quantifying the practical significance of differences (small ≈ 0.2, medium ≈ 0.5, large ≈ 0.8).

⚠ For non-parametric cases, Cliff's delta should also be computed.

```
1 strategies <- unique(df$gc_strategy)
2
3 # Generate all pairwise Cohen's d
4 pairwise_d <- combn(strategies, 2, simplify = FALSE, FUN = function(groups) {
5   d <- cohens_d(energy_joules ~ gc_strategy,
6                 data = df %>% filter(gc_strategy %in% groups))
7   d$comparison <- paste(groups, collapse = " vs ")
8   d
9 })
10
11 pairwise_d <- do.call(rbind, pairwise_d)
12 print(pairwise_d)
13
14 # Cliff's delta between groups (example GC1 vs GC2)
15 cliff.delta(energy_joules ~ gc_strategy, data = df)
16
17
```